



STANDARD OPERATING PROCEDURE

How and why we use AI-assisted tender analysis

Scope, rationale, and operating procedure for Pragati Opportunity Intelligence — prepared for executive review ahead of a pilot rollout to the Business Development team.

STATUS Prototype — internal review	OWNER Business Development & Technology
APPLIES TO Government & institutional tender intake	DATA CLASSIFICATION Public & synthetic data only

01 Purpose

This document explains what Pragati Opportunity Intelligence is, the business problem it was built to solve, how it fits into our existing tender process, and the procedure for using it. It is written for executive sign-off before the tool is used on live tender decisions.

02 Why we built this

Every government and institutional tender we receive — from the Army, CAPF, state police, BSF, and others — arrives as a dense 20–40 page document of technical specifications, eligibility rules, certification demands, and delivery terms. Today, deciding whether to pursue one requires a person to read the entire document, cross-reference it against what they know of our capabilities, and judge it by feel.

That process has three recurring costs to the business:

- **Speed.** A thorough read-through and capability check takes a trained analyst the better part of a day per tender, which limits how many opportunities we can properly evaluate.

- **Consistency.** Two analysts reading the same tender can reach different conclusions, because there is no shared, written-down standard for what "good fit" means.
- **Missed detail.** A single disqualifying clause on page 14 — a certification we don't hold, a turnover threshold we don't meet — is easy to miss under time pressure, and expensive to discover after a bid is already in motion.

Opportunity Intelligence was built to remove the first two costs entirely and sharply reduce the third, without removing the human decision at the end of the process.

03 Scope

IN SCOPE

- Reading and structuring any uploaded tender PDF
- Matching requirements against Pragati's *public* product & capability data
- Eligibility, certification, and compliance screening
- Risk flagging (delivery, capacity, competitive)
- One explainable score and a Pursue / Review / Do Not Pursue recommendation
- A full audit trail and mandatory human override capability

OUT OF SCOPE

- Any confidential or internal Pragati data
 - the knowledge base is public-source only
- Submitting or signing a bid — the system never acts on our behalf
- Legal or contractual review — this remains with Legal
- Final bid/no-bid authority — this stays with Business Development leadership

04 How it works

Every uploaded tender is passed through eight stages, run in a fixed order. The first six use AI to read and reason about the document; the last two are plain arithmetic and fixed business rules — not AI — so the final score and recommendation are always reproducible and auditable, not a black box.

- 1 **Document Intelligence**
Extracts clean text from the PDF, including scanned pages via OCR.
- 2 **Requirement Extraction**
Breaks the document into a structured list of individual requirements — not a summary, every clause.
- 3 **Capability Matching**
Checks each requirement against Pragati's public product data — Match, Partial Match, Gap, or Insufficient Evidence.
- 4 **Eligibility & Compliance**
Screens certification and eligibility clauses. Deliberately conservative — an unverified fact is never upgraded to a pass.
- 5 **Risk Analysis**
Flags delivery, capacity, and competitive risks, each with a mitigation note.
- 6 **Commercial & Strategic Analysis**
Assesses deal size and strategic fit against our current positioning.
- 7 **Opportunity Scoring** DETERMINISTIC
Combines the six factors above into one 0-100 score by fixed, published arithmetic — no AI judgment at this step.
- 8 **Decision** DETERMINISTIC
Applies fixed business rules to the score to produce Pursue / Review / Do Not Pursue.

WHY THE SPLIT MATTERS

AI is used where judgment is genuinely required — reading, matching, reasoning about risk. It is deliberately *not* used to set the final score or decision. That's fixed, auditable logic, so the same tender always produces the same score, and any factor of that score can be traced back to a specific page and a specific piece of evidence.

05 Operating procedure

The standard workflow for a Business Development analyst, from receipt of a tender to a recorded decision:

STEP	ACTION	OWNER
1	Upload the tender PDF to the system on receipt.	BD Analyst
2	Run analysis. The pipeline completes in minutes and produces a score, requirement list, and recommendation.	System
3	Review the Capability Match and Compliance tabs. Any item marked "insufficient evidence" is verified against internal records before proceeding.	BD Analyst
4	Review the Risk register and Commercial/Strategic notes for anything requiring escalation.	BD Analyst
5	Accept the AI recommendation, or override it with a written reason — every override is permanently logged.	BD Lead
6	Generate the Executive Decision Report and circulate for the standard bid/no-bid approval process.	BD Lead

This tool informs step 6 of our existing approval process — it does not replace it.

No bid decision is made without a person signing off.

06 What this is worth to the business

Hours → minutes First-pass read time per tender, freeing analyst time for the opportunities worth pursuing.	One standard Every tender is scored against the same six factors, not whoever happened to read it.	Zero silent gaps Every requirement is accounted for — nothing is skimmed past under deadline pressure.
---	--	--

The deeper value is trust, not just speed: because every score traces to cited evidence, leadership can question a recommendation and get a real answer — a page number and a quoted requirement — instead of "the AI said so."

07 Governance & data principles

- The capability knowledge base is built only from Pragati's own public-facing materials — no confidential or internal specifications are ingested.
- Compliance logic is conservative by design: a fact the system cannot verify is labelled "insufficient evidence," never treated as a pass or a fail.

- Every analysis, and every human override of it, is written to a permanent audit log with a timestamp and a reason.
 - A human can override the AI recommendation at any time; the AI cannot finalize or submit a bid.
-

08 Current status & limitations

This is a working prototype, not yet a production system, and should be read as such:

- It currently runs on public and clearly-labelled synthetic tender data for demonstration. Live tender use requires a short internal review before rollout.
 - The capability knowledge base reflects what's publicly published about Pragati today; it should be kept current as our product line evolves.
 - It is not connected to any internal system of record (ERP, certification database) — capability matches are only as complete as the public knowledge base behind them.
 - Every recommendation requires human review before it informs a bid decision. This is a design constraint, not a temporary limitation.
-

09 Recommended next steps

1. Approve a limited pilot: 5-10 live tenders, run in parallel with the existing manual process, to compare outcomes.
 2. Assign an owner to keep the capability knowledge base current as new products and certifications are published.
 3. Agree a review cadence (proposed: monthly) to check AI-flagged compliance items against ground truth and tune accuracy.
 4. Decide on a path to connect an internal capability data source, under proper access controls, once the pilot validates the approach.
-

10 Sign-off

Prepared by	Reviewed by — Business Development	Approved by — Chief Executive Officer
Name / Date	Name / Date	Name / Date

Pragati Opportunity Intelligence · Internal document · Uses public and synthetic data only · Not connected to any confidential system